

# Task 1 report:

## Pre-process the images :


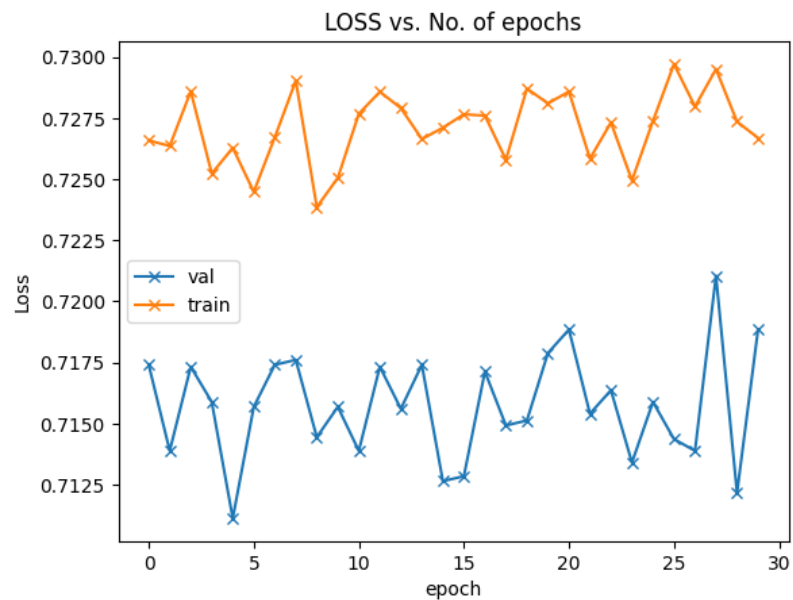
- 1) Resize to 128x128
- 2) Normalise the image: mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]

*Sample images after pre-process:*



**Model used** : Resnet50 with pretrain weights

Graphs:



The model's learning progress can be observed in the accuracy and loss graphs, but there appears to be minimal improvement in the model's performance. The accuracy graph shows a consistent range between 49 to 51%, and the loss graph follows a similar pattern. Despite the model's attempts to learn, there are no significant changes in its accuracy or loss values.

## **Results on evaluation step**

Accuracy: 0.496

Precision: 0.455

Recall: 0.14693

F1 score: 0.222

## Task 2 report:

Pre-process data :

```
[ ] 1 import re
    2
    3 # define a function to clean the text
    4 def clean_text(text):
    5     # remove unwanted characters such as punctuation marks, special characters, etc.
    6     cleaned_text = re.sub(r'^\w\s', '', text)
    7
    8     # print the removed characters
    9     removed_text = re.sub(r'\w\s', '', text)
   10     print(f"Removed characters: {removed_text}")
   11
   12     return cleaned_text
```

```
[ ] 1 # apply the clean_text function to the 'label' column of the pandas dataframe
    2 df_t['text'] = df_t['text'].apply(clean_text)
```

```
Removed characters:
Removed characters:
Removed characters: ..'
Removed characters: .
Removed characters:
Removed characters: +=
Removed characters: ::...
Removed characters:
Removed characters:
Removed characters:
Removed characters: -
Removed characters:
Removed characters:
Removed characters: ,,
```


### Model used :


Pretrained BERT model has been used in this case. To train the model for the above Hateful and Not-Hateful memes, last-layer fine tuning has been done.

Last Layer:

```
)  
(dropout): Dropout(p=0.1, inplace=False)  
(relu): ReLU()  
(fc1): Linear(in_features=768, out_features=512, bias=True)  
(fc2): Linear(in_features=512, out_features=2, bias=True)  
(softmax): LogSoftmax(dim=1)
```

Graphs:





## Result:

Accuracy: 0.594

Precision: 0.621

Recall: 0.43877

F1 score: 0.5143

```
Accuracy: 0.594
Classification Report:
              precision    recall  f1-score   support

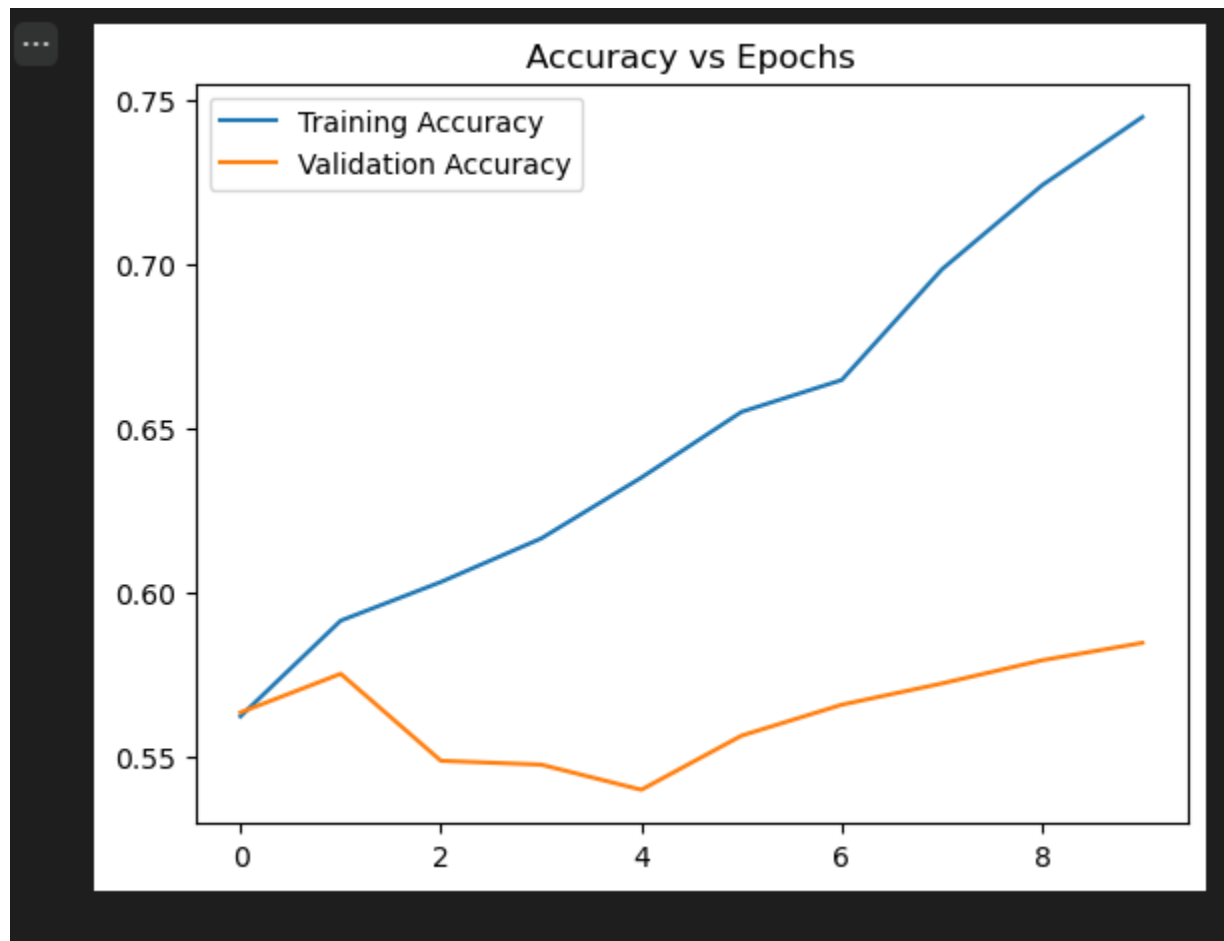
   Class 0       0.58      0.74      0.65        510
   Class 1       0.62      0.44      0.51        490

   accuracy          0.59          1000
  macro avg          0.60          1000
 weighted avg          0.60          1000
```


## Task 3 report:

### Graphs Of Loss and Accuracy

#### 1) Accuracy vs Epochs



## 2) Loss vs Epochs



## 3) Scores:


```
... Accuracy: 60.01%  
F1 score: 38.29%  
Precision: 52.78%  
Recall: 23.27%
```

### ClassWise Scores:


	precision	recall	f1-score	support
Non-hateful	0.63	0.77	0.70	1250
Hateful	0.40	0.25	0.31	750
accuracy			0.58	2000
macro avg	0.52	0.51	0.50	2000
weighted avg	0.55	0.58	0.55	2000

## 4) TSNE PLOTS

### TSNE for task 1:




### TSNE for Task 2:





### TSNE for Task 3:



1. From the TSNE plot of the TASK 1 model, it is evident that the model is unable to determine the hateful and non-hateful memes separately. Hence, more features are required for better classification.
2. From the TSNE plot of the TASK 2 model, it is evident that the model is unable to determine the hateful and non-hateful memes separately, given the fact that the model is scattering the data evenly. Hence, more features are required for better classification.
3. From the TSNE plot of the TASK 3 model, although not clearly, but two clusters are being formed from the use of the model. Hence the multimodal is performing better than the previous two models evidently.